

COS30018 – Task 2 Report

Name: Abhinav Karn

Student ID: 104488053

Date: 23rd August 2024

Explanation of the modifications and additions in v0.1 code as per instructed in Task 2:

1. Importing the Statements and Constants

```
1 import numpy as np
2 import matplotlib.pyplot as plt
3 import pandas as pd
4 import pandas_datareader as web
5 import datetime as dt
6 import tensorflow as tf
7 from sklearn.preprocessing import MinMaxScaler
8 from tensorflow.keras.models import Sequential
9 from tensorflow.keras.layers import Dense, Dropout, LSTM, InputLayer
10 import yfinance as yf
11 import os
```

These imports make it possible to use a number of very important libraries and functions, among them numpy and pandas for dealing with data, matplotlib for visualization of data, yfinance for fetching data, scikit-learn for scaling and splitting data, and finally, tensorflow for making a neural network. The 'os' module shall be used for operations dealing with files and directories, like checking for the existence of local data and for the creation of directories.

2. Defining constants and configuration parameters:

```
13 # Define constants and configuration parameters
14 N_STEPS = 50
15 LOOKUP_STEP = 1
16 TEST_SIZE = 0.2
17 FEATURE_COLUMNS = ['Adj Close', 'Volume']
18 LOSS = 'mean_squared_error'
19 UNITS = 256
20 N_LAYERS = 2
21 DROPOUT = 0.2
22 OPTIMIZER = 'adam'
23 BIDIRECTIONAL = False
24 SCALE = True
25 SPLIT_BY_DATE = True
26 SHUFFLE = False
27 MODEL_NAME = 'model_v0.1'
28 SAVE_LOCAL = True
29 DATA_DIR = 'local_data'
30 NAN_STRATEGY = 'fill_mean'
```

These constants can be thought of as parameters defining the configuration for data preprocessing and model training.

N_STEPS and LOOKUP_STEP control how big our input sequences are, and how far ahead we are predicting.

FEATURE_COLUMNS indicate the columns from the dataset used in making the predictions.

LOSS, UNITS, N_LAYERS, DROPOUT, OPTIMIZER, BIDIRECTIONAL These variables configure neural network architecture and training parameters.

It is determined by SCALE and NAN_STRATEGY whether to scale data or how to deal with NaN values.

SAVE_LOCAL, DATA_DIR, and MODEL_NAME manage local storage and naming conventions.

3. Loading the data and checking or the local storage

```
32 # Load data from Yahoo Finance
33 def load_data(ticker, start_date, end_date, feature_columns=FEATURE_COLUMNS):
34     local_file_path = f"{DATA_DIR}/{ticker}_data.csv"
35     if SAVE_LOCAL and os.path.exists(local_file_path):
36         data = pd.read_csv(local_file_path, index_col=0, parse_dates=True)
37         print(f"Loaded data from local storage: {local_file_path}")
38     else:
39         data = yf.download(ticker, start_date, end_date)
40         if SAVE_LOCAL:
41             os.makedirs(DATA_DIR, exist_ok=True)
42             data.to_csv(local_file_path)
43             print(f"Data saved locally to: {local_file_path}")
44
```

The code checks if the data for the given ticker is already in the local storage. If the data exists, then it directly reads it from a CSV file with `pd.read_csv()`. Otherwise, it downloads it with `yfinance` from Yahoo Finance. The `os.makedirs()` function makes sure to create the directory if it does not exist before saving the data, thus avoiding errors if the directory does not exist. This setup will give the best performance, reducing API calls as much as possible, especially during development.

4. Handling Missing Values Based on the Defined Strategy

```
48 # Handle missing values
49 if NAN_STRATEGY == 'drop':
50     data.dropna(inplace=True)
51 elif NAN_STRATEGY == 'fill_mean':
52     data.fillna(data.mean(), inplace=True)
53 else:
54     raise ValueError(f"Unknown NaN strategy: {NAN_STRATEGY}. Use 'drop' or 'fill_mean'.")
55
```

Explanation: This function handles missing values `NAN`. Here it applies the strategy defined in `NAN_STRATEGY`:

If it is set to `'drop'`, this removes all rows with missing values using `data.dropna(inplace=True)`.

If `'fill_mean'`, fills missing values with columns' mean using `data.fillna(data.mean(), inplace=True)`.

This flexibility will allow the model to process a dataset containing data of varying quality. The line `raise ValueError` ensures protection against strategy values that may be invalid.

5. Scaling Features Using MinMaxScaler

```
56 # Scale features
57 if SCALE:
58     scaler = MinMaxScaler(feature_range=(0, 1))
59     data[feature_columns] = scaler.fit_transform(data[feature_columns])
60 else:
61     scaler = None
62
```

It means that scaling of the features is important in a neural network, particularly when these features are of very different magnitude, such as stock prices compared to trading volumes. 'MinMaxScaler' will scale each feature onto a range of 0 through 1, normalizing the input data. As a result, it will make the LSTM model converge more

efficiently because all features are treated with equal importance. If set to 'False', no scaling takes place.

6. Splitting Data into Training and Testing Sets

```
63     # Split data
64     if SPLIT_BY_DATE:
65         train_data = data[:int(len(data) * (1 - TEST_SIZE))]
66         test_data = data[int(len(data) * (1 - TEST_SIZE)):]
67     else:
68         train_data, test_data = train_test_split(data, test_size=TEST_SIZE, shuffle=SHUFFLE, random_state=42)
69
```

The dataset is split according to the SPLIT_BY_DATE flag into a training and a test set:

In the event that SPLIT_BY_DATE is True, this will create a chronological split with a percentage defined by TEST_SIZE reserved for testing. This is typical for most time-series data where random shuffling might disrupt temporal dependencies.

Unless False, it uses train_test_split() from scikit-learn to randomly split the data. It has shuffling controlled by the SHUFFLE flag and a fixed random_state for reproducibility.

7. Creating Sequences for LSTM Input

```
# Create X and y arrays from the data
def create_xy(data, n_steps, lookup_step, feature_columns):
    """
    Create X (features) and y (labels) arrays from the data.

    Parameters:
    data (pd.DataFrame): Data containing features
    n_steps (int): Number of time steps to use as input
    lookup_step (int): Number of days to look ahead for the label
    feature_columns (list): List of feature columns to use

    Returns:
    X (numpy array): Array of feature data
    y (numpy array): Array of labels
    """
    X, y = [], []
    data = data[feature_columns].values
    for i in range(len(data) - n_steps - lookup_step):
        X.append(data[i:i+n_steps])
        y.append(data[i+n_steps+lookup_step-1][0]) # Use the first column for the label (e.g., 'adjclose')
    return np.array(X), np.array(y)
```

Overview: The following function is used to prepare the sequences of input (X) and output (y) while training the LSTM model:

The function processes the input by creating sub-arrays of length `n_steps`, which are past observations, in order to predict the value at `lookup_step` ahead.

This method of generating a sequence is how an LSTM learns to predict future values from the patterns in historical data.

The 0 index ensures this will only make use of the single column of features—like 'Adj Close'—for the prediction target.

8. Constructing the LSTM Model

```
98 # Construct the LSTM model
99 def create_model(n_steps, n_features, loss, units, n_layers, dropout, optimizer, bidirectional):
100     """
101     Create an LSTM model for time series prediction.
102
103     Parameters:
104     n_steps (int): Number of time steps in the input
105     n_features (int): Number of features in the input
106     loss (str): Loss function to use
107     units (int): Number of LSTM units in each layer
108     n_layers (int): Number of LSTM layers
109     dropout (float): Dropout rate to prevent overfitting
110     optimizer (str): Optimizer for training
111     bidirectional (bool): Whether to use bidirectional LSTM layers
112
113     Returns:
114     model (Sequential): Compiled LSTM model
115     """
116     model = Sequential()
117     for i in range(n_layers):
118         if i == 0:
119             if bidirectional:
120                 model.add(Bidirectional(LSTM(units, return_sequences=True), input_shape=(n_steps, n_features)))
121             else:
122                 model.add(LSTM(units, return_sequences=True, input_shape=(n_steps, n_features)))
123         elif i == n_layers - 1:
124             if bidirectional:
125                 model.add(Bidirectional(LSTM(units)))
126             else:
127                 model.add(LSTM(units))
128         else:
129             if bidirectional:
130                 model.add(Bidirectional(LSTM(units, return_sequences=True)))
131             else:
132                 model.add(LSTM(units, return_sequences=True))
133         model.add(Dropout(dropout))
134     model.add(Dense(1)) # Output layer for predicting a single value
135     model.compile(loss=loss, optimizer=optimizer)
136     return model
```

Overview: This function will construct the LSTM model.

Sequential() This is a linear stack of layers.

The loop adds LSTM layers. First and last layers are handled separately: input shape for the first, no return_sequences for the last to return a scalar prediction.

Bidirectional wrapping is conditional on the bidirectional flag and hence enables the model to learn patterns temporal in both directions.

The dropout layers prevent overfitting.

The model is compiled with the definition of a loss function and an optimizer to set what the goal for training shall be.

9. Inverse Scaling Predictions to Original Values

```
7 # Inverse transform to get actual price values (if scaling was applied)
8 if scaler:
9     predicted_prices = scaler.inverse_transform(np.concatenate((predicted_prices, np.zeros((predicted_prices.shape[0], len(FEATURE_COLUMNS) - 1))), axis=1))[:, 0]
10    true_prices = scaler.inverse_transform(np.concatenate((y_test.reshape(-1, 1), np.zeros((y_test.shape[0], len(FEATURE_COLUMNS) - 1))), axis=1))[:, 0]
11 else:
12    true_prices = y_test
```

This block rescales the scaled predictions and true values back to the original scale.

np.concatenate concatenates the prediction with zero-filled arrays to match the input shape expectation for a scaler.

scaler.inverse_transform This will inverse detach the scaled values back to the original stock prices so that the results are interpretable.

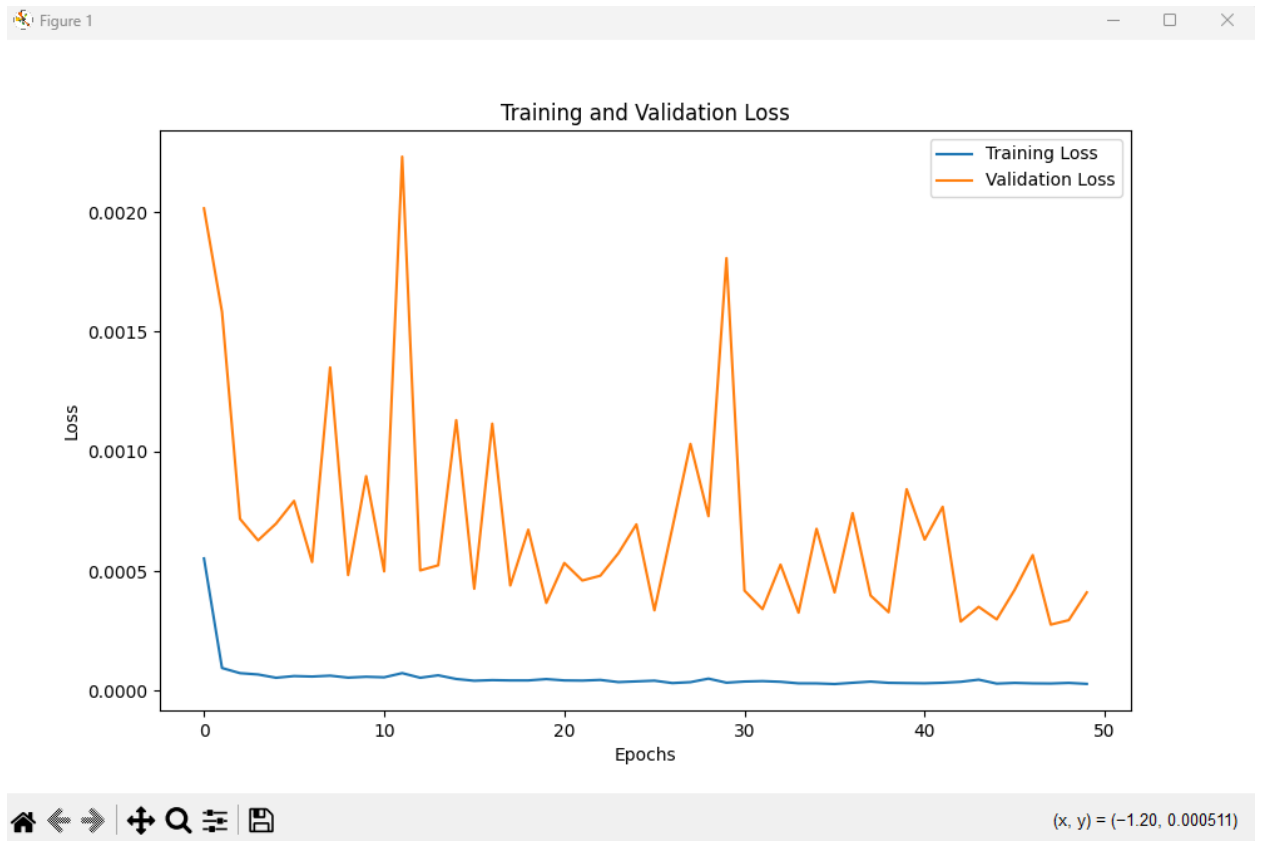
This is a crucial step because it ensures that the performance is evaluated and visualized accurately regarding real stock prices.

10. Plotting Actual vs. Predicted Prices

```
# Plot the true prices vs. predicted prices
plt.figure(figsize=(10, 6))
plt.plot(true_prices, color='black', label='Actual Prices')
plt.plot(predicted_prices, color='green', label='Predicted Prices')
plt.title('Actual vs. Predicted Prices')
plt.xlabel('Time')
plt.ylabel('Price')
plt.legend()
plt.show()
```

Explanation: This visualization plots actual stock prices against predicted prices:

- Helps visually assess the model's prediction performance.
- plt.plot() draws the lines for actual and predicted prices.
- The plot provides a clear picture of how well the model tracks real-world trends.



Summary:

In this detailed explanation, it shall cover all the important modifications and additions into my code.